Threaded Programming

Lecture 1: Concepts

Overview

- Shared memory systems

- Basic Concepts in Threaded Programming

Shared memory systems

- Threaded programming is most often used on shared memory parallel computers.

- A shared memory computer consists of a number of processing units (CPUs/cores) together with some memory

- Key feature of shared memory systems is a *single address space* across the whole memory system.
  - every CPU/core can read and write all memory locations in the system
  - one logical memory space
 - all CPUs/cores refer to a memory location using the same address

Conceptual model

Real hardware

- Real shared memory hardware is more complicated than this…..

  – Memory may be split into multiple smaller units

  – There may be multiple levels of cache memory

    • some of these levels may be shared between subsets of processors

  – The interconnect may have a more complex topology

- ….but a single address space is still supported

  – Hardware complexity can affect performance of programs, but not their correctness

Real hardware example

Threaded Programming Model

- The programming model for shared memory is based on the notion of threads
  - threads are like processes, except that threads can share memory with each other (as well as having private memory)
- Shared data can be accessed by all threads
- Private data can only be accessed by the owning thread
- Different threads can follow different flows of control through the same program
  - each thread has its own program counter
- Usually run one thread per CPU/core
  - but could be more
  - can have hardware support for multiple threads per core

Threads (cont.)

Thread 1 **Thread 2** **Thread 3**

| PC | Private data | | PC | Private data | | PC | Private data |

Shared data

Thread communication

- In order to have useful parallel programs, threads must be able to exchange data with each other

- Threads communicate with each other via reading and writing shared data
  - thread 1 writes a value to a shared variable A
  - thread 2 can then read the value from A

- Note: there is no notion of messages in this model

Thread Communication

	Thread 1	Thread 2

Program

`mya=23`
`a=mya`

`mya=a+1`

Private data

`23`

`24`

Shared data

`23`

Synchronisation

- By default, threads execute asynchronously

- Each thread proceeds through program instructions independently of other threads

- This means we need to ensure that actions on shared variables occur in the correct order: e.g.
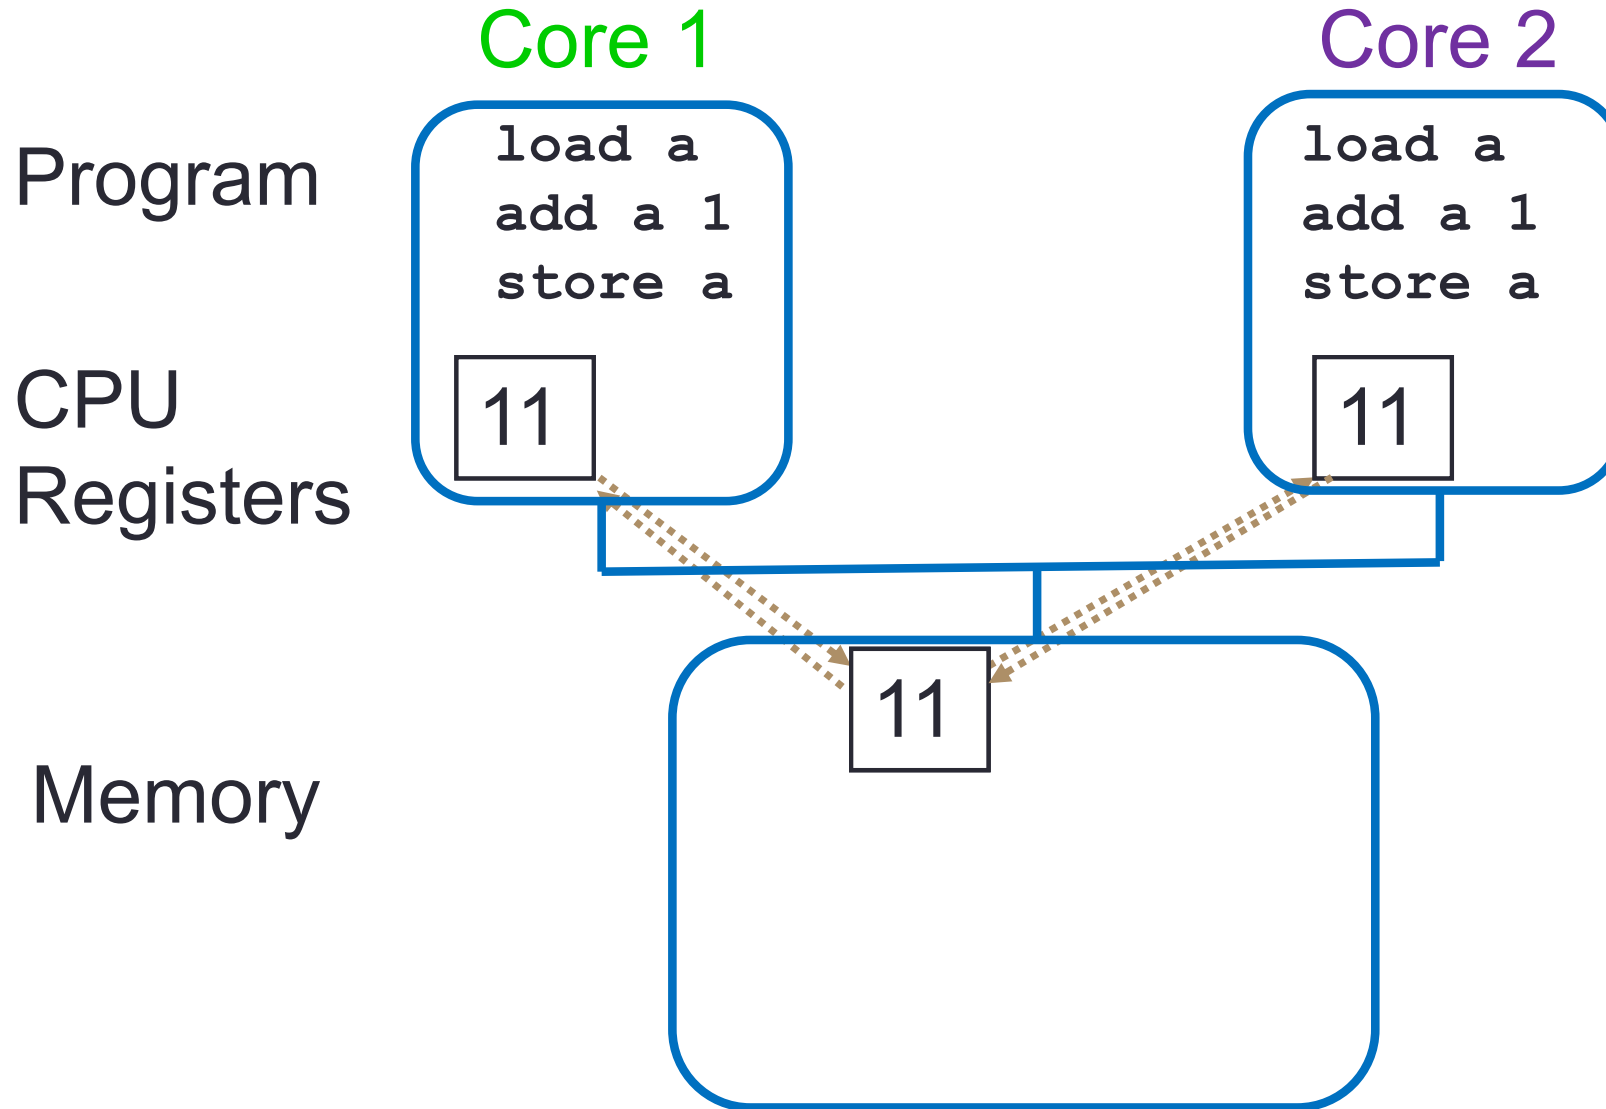
> thread 1 must write variable A before thread 2 reads it,
>
> or
>
> thread 1 must read variable A before thread 2 writes it.

- Note that updates to shared variables (e.g. `a = a + 1`) are *not* atomic!

- Requires (at least) 3 instructions: load, add, store

- If two threads try to do this "at the same time", the instructions can be interleaved in time such that one of the updates may get overwritten.

Synchronisation example

Core 1

Core 2

Program

```
load a
add a 1
store a
```

```
load a
add a 1
store a
```

CPU
Registers

11

11

Memory

11

|epcc|

Race conditions

- This problem is an example of a race condition.

- A race condition occurs when multiple threads access the same memory location without synchronization, and at least one of the accesses is a write.

- Without synchronization, different time orderings of the accesses may occur, possibly resulting in non-deterministic behavior

  - Running the same program with the same inputs twice may give different results

- Very nasty type of bug

  - Not reliably detectable by conventional software testing procedures

Tasks

- A *task* is a piece of computation which can be executed independently of other tasks

- In principle we could create a new thread to execute every task
  - in practise this can be too expensive, especially if we have large numbers of small tasks

- Instead tasks can be executed by a pre-exisiting *pool* of threads
  - tasks are submitted to the pool
  - some thread in the pool executes the task
  - at some point in the future the task is guaranteed to have completed

- Tasks may or may not be ordered with respect to other tasks

Parallel loops

- Loops are the main source of parallelism in many applications.

- If the iterations of a loop are *independent* (can be done in any order) then we can share out the iterations between different threads.

- e.g. if we have two threads and the loop

```
for (i=0; i<100; i++){
    a[i] += b[i];
}
```

  we could do iteration 0-49 on one thread and iterations 50-99 on the other.

- Can think of an iteration, or a set of iterations, as a task.

Reductions

- A *reduction* produces a single value from associative operations such as addition, multiplication, max, min, and, or.

- For example:

```
b = 0;

for (i=0; i<n; i++)
    b += a[i];
```

- Allowing only one thread at a time to update **b** would remove all parallelism.

- Instead, each thread can accumulate its own private copy, then these copies are reduced to give final result.

- If the number of operations is much larger than the number of threads, most of the operations can proceed in parallel